

Akhil Bodahanapati

747-312-1646 | akb198@pitt.edu | Pittsburgh, PA | [linkedin.com/in/akhil-bod](https://www.linkedin.com/in/akhil-bod) | github.com/akhilbod | akhilbod.github.io

EDUCATION

University of Pittsburgh

Bachelor of Science in Computer Science, GPA: 3.35

Pittsburgh, PA

Expected Dec 2027

- Relevant Coursework: Data Structures, Algorithms, Systems Software, Database Management, Machine Learning

EXPERIENCE

Nutanix

Software Engineer Intern

May 2026 – Aug 2026

San Jose, CA

- Redesigned partial-update logic to preserve cluster state, stopping concurrent updates from wiping customer configs
- Migrated per-customer rate limiting for the production IaaS control plane onto a new Go backend service
- Added concurrency checks rejecting conflicting writes to control-plane workflows on Kubernetes, Hyper-V, and RHEL
- Won the US + EMEA intern hackathon with a Kubernetes GPU scheduler that autoscaled and released idle GPUs

NLP and AI Lab, University of Pittsburgh

Software Engineer Intern

Dec 2025 – Apr 2026

Pittsburgh, PA

- Cut experiment turnaround from 4 min to 75 s (3.2x) by rewriting slow PostgreSQL queries and adding Redis job queues
- Built parallel Python inference pipelines over 1M+ documents, enabling corpus-scale NLP experiments for researchers
- Built async FastAPI services to submit, track, and retrieve long-running inference jobs for concurrent users

Federated Hermes

Software Engineer Intern

May 2025 – Aug 2025

Pittsburgh, PA

- Built Python/Flask services for portfolio analysis and rebalancing workflows on portfolios totaling \$2B+ in client assets
- Partnered with portfolio managers to replace manual portfolio and trade workflows, projected to save \$75K/yr in effort
- Designed PostgreSQL schemas linking portfolio, account, and trade records to power investment and reporting queries

PROJECTS

InterviewSense | interviewsense.org | 25,000+ Registered Users

- Built and scaled an AI interview platform with Python, FastAPI, Azure OpenAI, PostgreSQL, and React
- Designed Stripe billing and a usage-credit ledger that automatically refunds credits when model requests fail
- Engineered pipelines syncing 10,000+ job postings into search-indexed pages, driving 20K+ monthly organic visitors

Orbit | github.com/Not-Ethan/Orbit | CMU TartanHacks Winner

- Built a distributed semantic search engine over 500K records using Python, C++, and FAISS
- Partitioned vector search across independent workers and merged per-shard ranked results in a central coordinator

Quantitative Investment Engine | 1st Place, Federated Hermes Hackathon

- Built a Python and SQL research engine on Databricks and AWS that ranked investment ideas on historical market data
- Automated strategy backtesting and portfolio analysis into reproducible pipelines, replacing recurring manual research

SKILLS

Languages: Python, Go, C++, Java, TypeScript, JavaScript, SQL, Bash

Frameworks: FastAPI, Flask, React

Infrastructure: AWS, Docker, Kubernetes, Linux, Git, GitHub Actions

Data: PostgreSQL, Redis, Databricks, FAISS